

# Natural Language Processing using Python Programming

## Notebook 09.1: Introduction to Word Embeddings

Python 3.8+    NLTK Latest    SpaCy Latest    Gensim Latest    License MIT

Part of the comprehensive learning series: [Natural Language Processing using Python Programming](#)

### Learning Objectives:

- Understand the fundamental concepts and advantages of word embeddings over sparse vector representations
- Master the distributional hypothesis and its role in capturing semantic meaning
- Explore vector arithmetic and analogical reasoning capabilities of word embeddings
- Learn about popular embedding models including Word2Vec and GloVe architectures
- Build foundation for practical implementation of word embedding techniques

- This notebook introduces the revolutionary concept of **word embeddings** - dense vector representations that capture semantic meaning, moving beyond the limitations of sparse count-based methods like TF-IDF.
- We'll explore how embeddings solve the fundamental problem of semantic similarity and understand the distributional hypothesis that forms the foundation of modern NLP techniques.

## 1. The Problem with Sparse Vectors (Review)

**Sparse Vectors** (like TF-IDF) are typically very long (equal to the size of the vocabulary, often 100,000+ dimensions) and mostly filled with zeros. This is inefficient and ignores context.

Feature	TF-IDF (Sparse Vector)	Word Embedding (Dense Vector)
Dimensions	High (e.g., 50,000)	Low (e.g., 50, 100, 300)
Values	Integers or floating-point weights (mostly 0s)	Continuous floating-point numbers
Meaning	Measures word <i>importance</i> (count/rarity)	Measures word <i>meaning</i> (context/semantics)
Size	Grows with vocabulary size	Fixed size

## Why TF-IDF Fails at Meaning:

- If 'good' and 'excellent' never appear in the same document, TF-IDF will place them far apart, even though they mean similar things.
- They have no measurable similarity beyond being two different words.

## 2. Distributional Semantics: The Core Idea

- The foundation of all modern word embeddings is the **Distributional Hypothesis**:

**"You shall know a word by the company it keeps."** (Firth, 1957)

- If two words frequently appear in the same contexts (i.e., surrounded by the same words), they are likely to have similar meanings.
  - For example, 'cat' and 'dog' both frequently appear near 'owner', 'feed', and 'pet'.

## 3. Word Embeddings: How Semantic Space Works

- A word embedding is a multi-dimensional coordinate for a word.
- Words with similar meanings are mapped to positions close to each other in the vector space.

## The Analogy Task (Vector Arithmetic)

- The most famous illustration of embeddings is vector arithmetic, which shows that geometric relationships in the vector space correspond to semantic relationships in language:

$$\text{Vector}(\text{'King'}) - \text{Vector}(\text{'Man'}) + \text{Vector}(\text{'Woman'}) \approx \text{Vector}(\text{'Queen'})$$

- This works because the 'Royalty' dimension and the 'Gender' dimension are encoded into the dense vector space.

```
In [1]: # Conceptual visualization of a 2-D vector space
import matplotlib.pyplot as plt

words = ['King', 'Queen', 'Man', 'Woman', 'Apple', 'Banana']
x = [10, 10, 5, 5, -2, -3]
y = [15, 12, 10, 8, 3, 2]

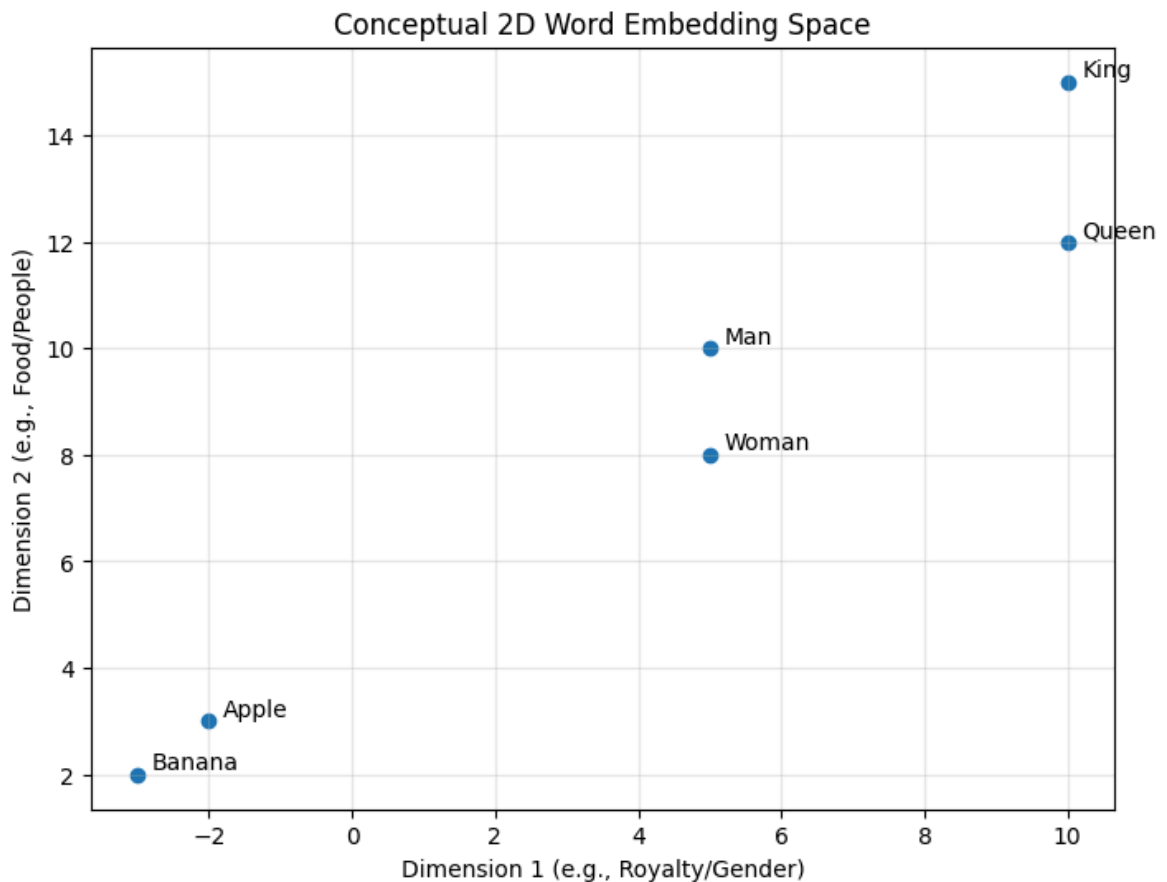
plt.figure(figsize=(8, 6))
plt.scatter(x, y)

for i, word in enumerate(words):
    plt.annotate(word, (x[i] + 0.2, y[i] + 0.1))

plt.title('Conceptual 2D Word Embedding Space')
plt.xlabel('Dimension 1 (e.g., Royalty/Gender)')
```

```
plt.ylabel('Dimension 2 (e.g., Food/People)')
plt.grid(True, alpha=0.3)
plt.show()
```

*# Observation: 'King' and 'Queen' are close. 'Apple' and 'Banana' are close.  
# The arithmetic relationship (King - Man + Woman) should land near Queen.*



## 4. Popular Embedding Models (Brief Overview)

### 4.1 Word2Vec (Mikolov, 2013)

- Developed by Google, Word2Vec uses a shallow neural network to learn word vectors from context.
- It has two main architectures:
  - **CBOW (Continuous Bag-of-Words):** Predicts the current word based on its surrounding context words.
  - **Skip-gram:** Predicts the surrounding context words given the current word (generally performs better).

### 4.2 GloVe (Global Vectors for Word Representation)

- Developed at Stanford, GloVe combines both global (corpus-wide) and local (contextual) matrix factorization information.
- It's often highly effective and is a popular choice for pre-trained vectors.

## 5. Summary and Next Steps

- Word embeddings revolutionize NLP by introducing dense, semantic-aware vector representations, fundamentally moving beyond count-based methods.
- In the next notebook (9.2), we will gain practical experience by working with **Gensim** to load and train **Word2Vec** models, and visualize these semantic spaces using **t-SNE**.

### Key Takeaways

- **Semantic Revolution:** We learned how word embeddings revolutionize NLP by capturing semantic meaning through dense vector representations, solving the fundamental limitations of sparse count-based methods.
  - **Distributional Hypothesis Mastery:** We understood that "words are known by the company they keep" - the core principle that enables embeddings to capture contextual similarity and meaning.
  - **Vector Arithmetic Wonder:** We explored the fascinating capability of embeddings to perform analogical reasoning through vector arithmetic (King - Man + Woman  $\approx$  Queen).
  - **Model Architecture Understanding:** We gained insight into popular embedding models like Word2Vec (CBOW/Skip-gram) and GloVe, understanding their different approaches to learning semantic representations.
- 

### Next Notebook Preview

- With theoretical foundations established, we're ready for **hands-on implementation** of word embeddings.
  - Notebook 9.2 will provide practical experience with **Gensim** for loading and training Word2Vec models, plus visualization techniques using t-SNE to explore semantic spaces.
- 

### About This Project

This notebook is part of the **Natural Language Processing using Python Programming for Beginners** repository - a comprehensive, beginner-friendly guide for mastering NLP using Python, NLTK, and SpaCy.

**Repository:** `NLP`

**Author**

**Prakash Ukhalkar**



---

Built with ❤️ for the Python community