

Natural Language Processing using Python Programming

Notebook 10.1: Introduction to Transformer Models (BERT, GPT)

Python3.8+

TransformersLatest

PyTorchLatest

PandasLatest

LicenseMIT

Part of the comprehensive learning series: [Natural Language Processing using Python Programming](#)

Learning Objectives:

- Understand the revolutionary Transformer architecture and self-attention mechanisms
- Master the Hugging Face ecosystem and pipeline utilities for instant NLP applications
- Implement zero-shot classification without training data for new categories
- Explore advanced NLP tasks including NER, question answering, and text summarization
- Build foundation for fine-tuning pre-trained models on custom datasets

- This notebook introduces the revolutionary **Transformer architecture** that transformed modern NLP, featuring self-attention mechanisms that enable models to understand context better than ever before.
- We'll explore the **Hugging Face ecosystem** and learn to leverage pre-trained models like BERT and GPT through simple pipeline utilities for immediate, powerful NLP applications without complex setup.

1. The Transformer Architecture (Conceptual)

The original Transformer has two main blocks: the **Encoder** and the **Decoder**. Modern models typically use only one:

Model Type	Primary Function	Key Example	Typical Tasks
Encoder-Only	Understanding (Reads the entire input sentence)	BERT	Classification, NER, Feature Extraction
Decoder-Only	Generation (Generates text sequentially)	GPT	Translation, Chatbots, Text Completion

The Key Idea: Self-Attention

Self-attention allows the model to resolve ambiguity by determining which words in a sentence are relevant to a specific word.

Example: In the sentence, "The animal didn't cross the road because it was too wide." The model learns that 'it' refers to 'the road' by placing high attention weight on that word.

2. Introducing Hugging Face and the pipeline

The **Hugging Face transformers** library provides a standardized way to access hundreds of pre-trained models. The easiest way to start is with the **pipeline** utility, which handles all preprocessing, model loading, and post-processing in one line.

```
In [1]: # Suppress specific user warnings for transformers library
import os
os.environ["HF_HUB_DISABLE_SYMLINKS_WARNING"] = "1"
```

```
In [4]: import warnings

# Suppress TF deprecation warnings
warnings.filterwarnings("ignore", category=DeprecationWarning)
```

```
In [5]: # Import necessary libraries
# transformers library is used for working with transformer models
# pipeline is a high-level API for various NLP tasks
from transformers import pipeline
import pandas as pd

print("Hugging Face Transformers library loaded.")
```

Hugging Face Transformers library loaded.

2.1 Practical Application: Zero-Shot Classification

Zero-Shot Classification is the ability to classify a document into categories **without any training data** for those specific categories.

```
In [6]: # The Zero-Shot pipeline loads a suitable pre-trained model (usually BART)
# BART: Bidirectional and Auto-Regressive Transformers
# BART is particularly effective for text generation and comprehension tasks
# We use the "facebook/bart-large-mnli" model for zero-shot classification
# which is fine-tuned for natural language inference (NLI)
classifier = pipeline("zero-shot-classification", model="facebook/bart-large-mnli")

text = "The quarterly earnings report showed massive gains in the tech sector, c
candidate_labels = ["finance", "sports", "politics", "technology"]

# Run the classification
result = classifier(text, candidate_labels)
```

```
print("Zero-Shot Classification Results:")
# Use Pandas to display results neatly
df_result = pd.DataFrame({'label': result['labels'], 'score': result['scores']})
print(df_result.sort_values(by='score', ascending=False).reset_index(drop=True))
```

Zero-Shot Classification Results:

	label	score
0	technology	0.9895
1	finance	0.0047
2	sports	0.0036
3	politics	0.0022

3. Other Transformer Applications via Pipeline

The `pipeline` is versatile and supports many complex NLP tasks instantly.

3.1 Named Entity Recognition (NER)

```
In [9]: # NER pipeline (often uses a BERT-based model)
# BERT: Bidirectional Encoder Representations from Transformers
# BERT is designed to understand the context of words in a sentence
# We use the default NER model provided by the transformers library
from transformers import logging
# suppress warnings
logging.set_verbosity_error()

ner_pipeline = pipeline(
    "ner",
    model="dbmdz/bert-large-cased-finetuned-conll03-english",
    revision="4c53496",
    grouped_entities=True
)

context = "Jeff Bezos founded Amazon in Seattle, Washington in 1994, and he retired as CEO in 2021."

print(f"\nContext: {context}")
print("\nExtracted Entities:")
ner_results = ner_pipeline(context)

for entity in ner_results:
    print(f" - Entity: {entity['word'][:15] | Type: {entity['entity_group'][:5] | Score: {entity['score'][:10]})")
```

Context: Jeff Bezos founded Amazon in Seattle, Washington in 1994, and he retired as CEO in 2021.

Extracted Entities:

- Entity: Jeff Bezos	Type: PER	Score: 0.9748
- Entity: Amazon	Type: ORG	Score: 0.9984
- Entity: Seattle	Type: LOC	Score: 0.9985
- Entity: Washington	Type: LOC	Score: 0.9984

3.2 Question Answering (Q&A)

```
In [12]: qa_pipeline = pipeline("question-answering")

context = "The new data center was built in Frankfurt, Germany in 2023. " \
```

```
"The total investment was $50 million, and it uses 100% renewable energy."
```

```
question1 = "How much did the new data center cost?"
question2 = "Where is the new data center located?"
question3 = "When was the new data center built?"

answer1 = qa_pipeline(question=question1, context=context)
answer2 = qa_pipeline(question=question2, context=context)
answer3 = qa_pipeline(question=question3, context=context)
print(f"\nContext: {context}")
print(f"\nQuestion 1: {question1}")
print(f"Answer: {answer1['answer']} (Score: {answer1['score']:.4f})")
print(f"\nQuestion 2: {question2}")
print(f"Answer: {answer2['answer']} (Score: {answer2['score']:.4f})")
print(f"\nQuestion 3: {question3}")
print(f"Answer: {answer3['answer']} (Score: {answer3['score']:.4f})")
```

Context: The new data center was built in Frankfurt, Germany in 2023. The total investment was \$50 million, and it uses 100% renewable energy.

Question 1: How much did the new data center cost?
Answer: \$50 million (Score: 0.9436)

Question 2: Where is the new data center located?
Answer: Frankfurt, Germany (Score: 0.9495)

Question 3: When was the new data center built?
Answer: 2023 (Score: 0.9624)

3.3 Text Summarization

```
In [14]: # Summarization pipeline (often uses models like T5 or BART)
summarizer = pipeline("summarization")

long_text = """
Recent advancements in artificial intelligence have brought Large Language Models like GPT-4 and Claude to the forefront of technology. These models, built upon the Transformer architecture, are capable of tasks ranging from creative writing to complex code generation. While they consume enormous computational resources for training, their widespread application is now democratizing access to advanced reasoning capabilities across various industries, from healthcare to finance.
"""

print(f"\nOriginal Text Length: {len(long_text)} characters")
# Note: Running this may require a model download.
summary = summarizer(long_text, max_length=50, min_length=10, do_sample=False)[0]

print(f"Summary: {summary}")
```

Original Text Length: 495 characters

Summary: Large Language Models (LLMs) built upon the Transformer architecture are capable of tasks ranging from creative writing to complex code generation. While they consume enormous computational resources for training, their widespread application is now democratizing access

4. Summary and Next Steps

- We introduced the core concepts of the **Transformer architecture** and demonstrated the immediate, powerful utility of the **Hugging Face pipeline**

for high-value tasks like Zero-Shot Classification, NER, and Q&A.

- This showcases the incredible leap in NLP capability.
- The next step is to adapt this power to your own data via **fine-tuning**.

Key Takeaways

- **Transformer Revolution:** We learned how the Transformer architecture revolutionized NLP through self-attention mechanisms, enabling better context understanding than previous RNN/LSTM models.
- **Hugging Face Mastery:** We mastered the Hugging Face ecosystem and pipeline utilities, providing instant access to state-of-the-art models with minimal code complexity.
- **Zero-Shot Capabilities:** We discovered the power of zero-shot classification, allowing document categorization without training data for specific categories.
- **Multi-Task Applications:** We explored diverse NLP applications including NER, question answering, and text summarization, demonstrating the versatility of Transformer models.

Next Notebook Preview

- With Transformer fundamentals established, we're ready to **customize these powerful models** for specific use cases.
- Notebook 10.2 will guide you through **fine-tuning BERT** for classification tasks, adapting pre-trained models to your own datasets and requirements.

About This Project

This notebook is part of the **Natural Language Processing using Python Programming for Beginners** repository - a comprehensive, beginner-friendly guide for mastering NLP using Python, NLTK, and SpaCy.

Repository: [NLP](#)

Author

Prakash Ukhalkar

